

CustomImageDataset class is doing.

Summary of CustomImageDataset

1. Initialization (__init__)

- Takes the dataset folder (root_dir) and optional transforms.
- Finds all **class folders** (ignores hidden ones).
- Creates a mapping:

```
{"class1": 0, "class2": 1, ...}
```

- Calls _load_images() to collect all image file paths and their class labels.
 - Stores them in self.images .
-

2. Image Collection (_load_images)

- Scans each class folder inside root_dir .
- For every valid image file (.png , .jpg , .jpeg , ...), saves:

```
(image_path, class_index)
```

- Example:

```
[("root_dir/class1/img1.jpg", 0),  
 ("root_dir/class1/img2.jpg", 0),  
 ("root_dir/class2/img3.jpg", 1)]
```

3. Length (__len__)

- Returns the total number of images found.
 - Used by len(dataset) and DataLoader.
-

4. Get Item (__getitem__)

- Given an index idx :
 - Looks up (img_path, label) from self.images .
 - Opens the image and **converts it to grayscale ('L') → 1 channel**.
 - Applies transforms if provided.

- Returns (image, label).

✓ In short:

This class:

1. Scans your dataset directory once,
2. Maps each image to its class label,
3. Loads images as **grayscale**,
4. Applies transforms,
5. Returns (image_tensor, label) when accessed.

👉 The Dataset base class in PyTorch is **just an interface** — it only defines that your custom class must implement:

- `__len__(self)`
- `__getitem__(self, idx)`

That's it. It does **not** handle class detection or label mapping for you.

🔍 Where does mapping happen then?

- In your code, **you wrote it manually** here:

```
self.classes = [d for d in os.listdir(root_dir) if os.path.isdir(os.path.join(root_dir, d))]
self.classes.sort()
self.class_to_idx = {cls_name: i for i, cls_name in enumerate(self.classes)}
```

So:

- `self.classes = list of folder names (["cats", "dogs"])`
- `self.class_to_idx = mapping ({"cats": 0, "dogs": 1})`

Later, when `_load_images` runs, you use this mapping to assign each image a label.

⚡ But in `torchvision.datasets.ImageFolder`

Yes — the mapping *does* happen automatically.

- `ImageFolder` assumes your dataset is structured like:

```
root/
├── class1/
│   └── img1.jpg
```

```
└─ class2/
    └─ img2.jpg
```

- It automatically:
 - Detects the class folders.
 - Sorts them.
 - Creates `class_to_idx`.
 - Returns `(image_tensor, label)` when you access an item.

So your `CustomImageDataset` is basically **re-implementing what `ImageFolder` already does**, but with your own tweaks (like grayscale conversion).

✅ **Answer in one line:** No, the mapping doesn't happen automatically in a custom `Dataset` — you must implement it (as you did). But if you use `torchvision.datasets.ImageFolder`, then yes, mapping happens automatically.

Do you want me to show you how your class could be simplified just by using **`ImageFolder` + a small transform** (instead of writing `_load_images` etc.)?

✓ ResNet50

1. Mount Google Drive

- Mounts Google Drive to access dataset stored in [/content/drive/MyDrive/Project2/Project 1 Data](#).

2. Device Configuration

- Checks for GPU availability (`cuda`) and sets the device for training.

3. Data Augmentation and Normalization

- **Basic Transform:** Resize, center crop, convert to tensor, normalize using ImageNet stats.
- **Augmentation Transform:** Random horizontal flip, rotation, color jitter, affine transforms for creating diverse samples.

4. Load and Split Dataset

- Loads images using `ImageFolder`.
- Extracts labels for stratified splitting.

- Splits the dataset into **80% training** and **20% validation** using `StratifiedShuffleSplit`.
-

5. Dataset Balancing

- Computes class counts for the training subset.
 - Determines maximum class count and calculates **augmentation needed per class**.
 - Defines a **custom `BalancedDataset`** class:
 - Original samples use `basic_transform`.
 - Additional samples are augmented using `augment_transform + basic_transform`.
 - Verifies that the training dataset is now balanced.
-

6. Prepare Validation Dataset

- Defines a `SubsetWithTransform` class to apply `basic_transform` to validation data.
-

7. Create DataLoaders

- Creates **`train_loader`** and **`val_loader`** with batch size 32.
 - Shuffles training data, does not shuffle validation data.
 - Uses multiple workers (`num_workers=4`) for faster data loading.
-

8. Model Setup (Transfer Learning)

- Loads a pre-trained **`ResNet50`** model using `timm`.
 - Replaces the final fully connected layer to match **5 classes**.
 - Moves model to GPU or CPU.
-

9. Loss Function and Optimizer

- Uses **cross-entropy loss** for multi-class classification.
 - Uses **Adam optimizer** with learning rate `3e-5` and L2 regularization (`weight_decay=1e-5`).
 - Adds a **StepLR scheduler** to decrease learning rate every 7 epochs.
-

10. Training and Validation Function

- Loops over `num_epochs` :
 - **Training Phase:** forward pass, compute loss, backpropagate, update weights.
 - **Validation Phase:** compute validation loss and accuracy without gradient.
- Tracks **best validation accuracy** and saves the model if it improves.

- Prints epoch-wise statistics: train loss, validation loss, validation accuracy, and time.

11. Model Saving

- Saves the **best performing ResNet model** as `best_res_net_model.pth`.

✅ **In short:** You are **loading, augmenting, and balancing a dataset**, training a **pre-trained ResNet50** with a custom classifier, tracking validation performance, and saving the best model.

✓ ViT base Model

Here's a **pointwise summary** of what your code is doing:

1. Device Configuration:

- Checks if GPU is available and sets the device to `cuda` or `cpu`.

2. Data Augmentation and Normalization:

- Defines `basic_transform` for resizing, cropping, tensor conversion, and normalization.
- Defines `augment_transform` for image augmentation (flip, rotation, color jitter, affine transform).

3. Load and Split Dataset:

- Loads images from Google Drive using `ImageFolder`.
- Splits dataset into 80% training and 20% validation using stratified shuffle split.

4. Class Balancing:

- Calculates class counts in training subset.
- Determines how many augmented samples are needed per class to balance dataset.
- Defines `BalancedDataset` class to generate augmented samples for underrepresented classes.
- Creates a balanced training dataset.

5. Prepare Validation Dataset:

- Wraps validation subset in `SubsetWithTransform` class to apply `basic_transform`.

6. Create DataLoaders:

- Defines `DataLoader` for training (`shuffle=True`) and validation (`shuffle=False`) with batch size 32 and 4 workers.

7. Model Setup (Transfer Learning - ViT):

- Loads pre-trained Vision Transformer (`vit_base_patch16_224`).
- Replaces the classification head with a new linear layer to match `num_classes = 5`.
- Moves the model to the configured device.

8. Loss Function, Optimizer, Scheduler:

- Uses `CrossEntropyLoss`.
- Uses Adam optimizer with weight decay for L2 regularization.
- Uses `StepLR` scheduler to decay learning rate every 7 epochs.

9. Training and Validation Function (`train_model`):

- Loops over epochs:
 - **Training Phase:** Sets model to train mode, computes loss, backpropagates, and updates weights.
 - **Validation Phase:** Sets model to eval mode, computes loss and accuracy.
- Tracks best validation accuracy and saves model whenever it improves.
- Prints epoch-wise train loss, validation loss, validation accuracy, and epoch time.

10. Model Saving:

- The best model (with highest validation accuracy) is saved as `'best_vit_model.pth'`.

✓ Vit FineTune

✓ 1. Device Setup

- Using GPU if available (`cuda`), otherwise CPU.

2. Data Preparation

- **Basic transforms:** Resize to 256, center crop to 224, convert to tensor, normalize with ImageNet stats.
- **Augmentation:** Random horizontal flip, rotation, color jitter, affine transforms to generate additional samples.

3. Dataset Split

- Training and validation sets are created with a **balanced validation set** (~10% of data).
- `Subset` and custom dataset classes are used to handle transforms and maintain class balance.

4. Class Balancing

- Compute class counts in training set.
- Determine how many augmented samples are needed for each class to match the **maximum class count**.
- `BalancedDataset` applies augmentation to under-represented classes while keeping original samples.

5. DataLoaders

- Training and validation DataLoaders created with batch size 32 and multiple workers.

6. Model Setup (ViT)

- Pre-trained ViT model (`vit_base_patch16_224`) is loaded.
- **Classification head replaced** for 5 classes with added dropout:

```
model.head = nn.Sequential(  
    nn.Dropout(p=0.5),  
    nn.Linear(model.head.in_features, num_classes)  
)
```

- **Fine-tuning type:**
 - **Full fine-tuning:** All layers, including pre-trained ones, are updated during training because all parameters are passed to the optimizer.
 - Classification head is initialized randomly, learning from scratch.

7. Loss, Optimizer, Scheduler

- Loss: `CrossEntropyLoss`.
- Optimizer: Adam with learning rate $3e-5$ and L2 regularization.
- Learning rate scheduler: `StepLR` (reduces LR every 7 epochs by 0.1 factor).

8. Training and Validation

- **Training loop:**
 - Model set to train mode.
 - Forward pass → compute loss → backward → optimizer step.
- **Validation loop:**
 - Model set to eval mode.
 - Compute loss and accuracy without gradient computation.

- Best model saved based on **highest validation accuracy**.

9. Output

- Trains for 20 epochs.
 - Saves the best model (`best_vit_model.pth`).
 - Reports **best validation accuracy** at the end.
-

✓ In short:

- Full ViT model is fine-tuned on a small, balanced, and augmented dataset with a replaced classification head for 5 classes.
 - Augmentation and balancing are used to improve generalization.
 - All layers, including pre-trained ones, are updated, so this is **full fine-tuning** rather than just head training.
-

Start coding or [generate](#) with AI.